

EX 2 MULTI-LAYER PERCEPTRON (MLP) FOR CLASSIFICATION

DATE:

Problem Statement:

Develop a Multi-Layer Perceptron (MLP) for a simple classification task. Experiment with different numbers of hidden layers and activation functions, and evaluate the model's performance using accuracy and loss.

Suggested Dataset: Iris Dataset

Objectives:

1. Understand the structure and purpose of MLPs for classification.
2. Experiment with various hidden layer configurations and activation functions.
3. Train the MLP using the Iris dataset and evaluate its accuracy and loss.
4. Visualize training progress and use the trained model for prediction.

Scope:

This experiment provides insights into how neural networks with multiple layers (MLPs) perform on structured classification tasks. It demonstrates model tuning using different architectures and activation functions, an essential concept in designing effective deep learning models.

Tools and Libraries Used:

1. Python 3.x
2. TensorFlow / Keras
3. scikit-learn (for data preprocessing and dataset loading)
4. Matplotlib (for visualization)

Implementation Steps:

Step 1: Import Necessary Libraries

```
import matplotlib.pyplot as plt
import numpy as np
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

Step 2: Load and Preprocess Data

```
iris = load_iris()
X = iris.data
y = iris.target.reshape(-1, 1)
```

```
classes = iris.target_names
```

```
encoder = OneHotEncoder(sparse_output=False)
```

```
y_encoded = encoder.fit_transform(y)
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_encoded, test_size=0.2,  
random_state=42)
```

Step 3: Define MLP Model Creation Function

```
def create_mlp(input_dim, output_dim, hidden_layers, activation='relu'):
```

```
    model = Sequential()
```

```
    model.add(Dense(hidden_layers[0], input_dim=input_dim, activation=activation))
```

```
    for units in hidden_layers[1:]:
```

```
        model.add(Dense(units, activation=activation))
```

```
    model.add(Dense(output_dim, activation='softmax'))
```

```
    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
    return model
```

Step 4: Train and Evaluate MLP with Various Configurations

```
hidden_layer_configs = [[8], [16, 8], [32, 16, 8]]
```

```
activations = ['relu', 'tanh', 'sigmoid']
```

```
for hidden_layers in hidden_layer_configs:
```

```
    for activation in activations:
```

```
        print(f"\nTesting MLP with hidden_layers={hidden_layers}, activation={activation}")
```

```
        model = create_mlp(input_dim=4, output_dim=3, hidden_layers=hidden_layers,
```

```
activation=activation)
```

```
        history = model.fit(X_train, y_train, epochs=50, batch_size=5, verbose=0,  
validation_split=0.1)
```

```
        test_loss, test_acc = model.evaluate(X_test, y_test, verbose=0)
```

```
        print(f"Test Accuracy: {test_acc:.4f}, Test Loss: {test_loss:.4f}")
```

Step 5: Plot Accuracy and Loss Curves

```
    # Plot accuracy and loss
```

```
    plt.figure(figsize=(10, 4))
```

```
    plt.suptitle(f"Config: {hidden_layers}, Activation: {activation}", fontsize=14)
```

```
    plt.subplot(1, 2, 1)
```

```
    plt.plot(history.history['accuracy'], label='Train Acc')
```

```
    plt.plot(history.history['val_accuracy'], label='Val Acc')
```

```
    plt.xlabel('Epoch')
```

```
    plt.ylabel('Accuracy')
```

```

plt.title('Model Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Model Loss')
plt.legend()

plt.tight_layout()
plt.show()

```

Step 6: Predict on New Input

```

sepal_length = float(input("Sepal length (cm): "))
sepal_width = float(input("Sepal width (cm): "))
petal_length = float(input("Petal length (cm): "))
petal_width = float(input("Petal width (cm): "))

user_input = np.array([[sepal_length, sepal_width, petal_length, petal_width]])
user_input_scaled = scaler.transform(user_input)


# Predict using the last trained model
prediction = model.predict(user_input_scaled)
predicted_class_index = np.argmax(prediction)
predicted_class_name = classes[predicted_class_index]

print(f"\n 🌸 Predicted Iris Species: {predicted_class_name}")

```


```
Testing MLP with hidden_layers=[8], activation=relu
/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/dense.py:106: UserWarning: Do not pass an "input_shape"/"input_dim" argument to a layer. When using Sequential models, prefer using an "Input(shape)" object as the first layer in the model.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Test Accuracy: 0.9667, Test Loss: 0.1872
```

Config: [8], Activation: relu




```
Testing MLP with hidden_layers=[8], activation=tanh
Test Accuracy: 0.9667, Test Loss: 0.2080
```

Config: [8], Activation: tanh




Testing MLP with hidden_layers=[8], activation=sigmoid
Test Accuracy: 0.9333, Test Loss: 0.3909

Config: [8], Activation: sigmoid





Testing MLP with hidden_layers=[16, 8], activation=relu
Test Accuracy: 1.0000, Test Loss: 0.0655

Config: [16, 8], Activation: relu





Testing MLP with hidden_layers=[16, 8], activation=tanh
Test Accuracy: 1.0000, Test Loss: 0.0806

Config: [16, 8], Activation: tanh





Testing MLP with hidden_layers=[16, 8], activation=sigmoid
Test Accuracy: 0.9667, Test Loss: 0.3509

Config: [16, 8], Activation: sigmoid





Testing MLP with hidden_layers=[32, 16, 8], activation=relu
Test Accuracy: 1.0000, Test Loss: 0.0417

Config: [32, 16, 8], Activation: relu




Testing MLP with hidden_layers=[32, 16, 8], activation=tanh
Test Accuracy: 1.0000, Test Loss: 0.0429

Config: [32, 16, 8], Activation: tanh




Testing MLP with hidden_layers=[32, 16, 8], activation=tanh
Test Accuracy: 1.0000, Test Loss: 0.0429

Config: [32, 16, 8], Activation: tanh



Testing MLP with hidden_layers=[32, 16, 8], activation=sigmoid
Test Accuracy: 0.9333, Test Loss: 0.3200

Config: [32, 16, 8], Activation: sigmoid



Sepal length (cm): 10
Sepal width (cm): 5
Petal length (cm): 20
Petal width (cm): 15
1/1 — 0s 87ms/step

🌸 Predicted Iris Species: virginica

Conclusion:

This experiment demonstrates the use of MLPs for a multi-class classification problem using the Iris dataset. Through experimentation with different hidden layer depths and activation functions, students gain a practical understanding of neural network design and optimization techniques.